


Taming OpenStack with Ansible

OpenStack Juno Summit - May 2014



John Dewey



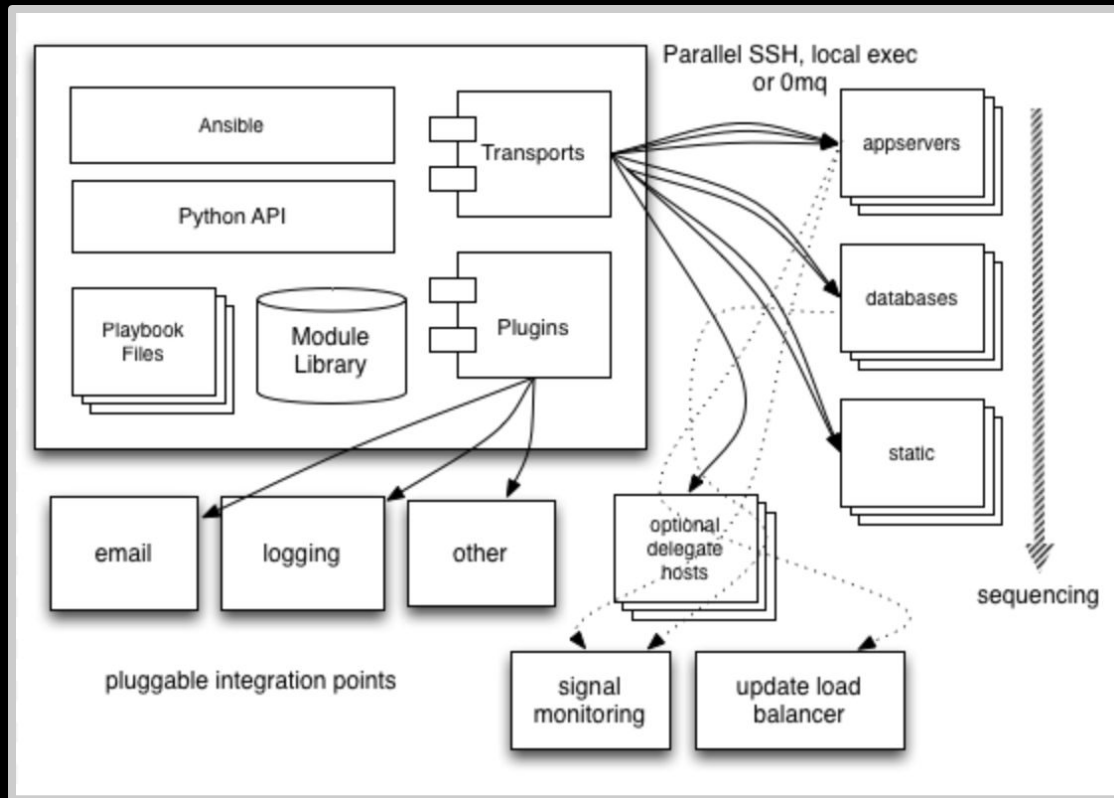
@retrOh

What is Ansible?

What is Ansible?

- Simple, multi-tier application deployment
- Deploys reliably and consistently
- Uses simple task descriptions vs custom code
- Capable of task orchestration

What is Ansible?



Overview

1. Simplicity in mind

Overview

1. Simplicity in mind
 - a. Easy to Install

- apt
- yum
- pip
- git clone
- tar

Requirements

- python 2.4+
- ssh

Overview

1. Simplicity in mind
 - a. Easy to Install
 - b. Serverless



Overview

1. Simplicity in mind
 - a. Easy to Install
 - b. Serverless
 - c. Agentless *



Overview

1. Simplicity in mind
 - a. Easy to Install
 - b. Serverless
 - c. Agentless *
 - d. YAML*

- name: add packages
- apt: pkg={{ item }}
- with_items:
 - smem
 - socat
 - pstack

Overview

1. Simplicity in mind
 - a. Easy to Install
 - b. Serverless
 - c. Agentless *
 - d. YAML*

Loops

- `with_items`
- `with_nested`
- `with_dict`
- `with_fileglob`
- `with_together`
- `with_subelements`
- `with_sequence`
- `with_random_choice`
- `with_first_found`
- `with_lines`
- `with_indexed_items`
- `with_flattened`

Overview

1. Simplicity in mind
 - a. Easy to Install
 - b. Serverless
 - c. Agentless *
 - d. YAML*

Conditionals

```
---
tasks:
  - command: /bin/false
    register: result
    ignore_errors: True
  - command: /bin/something
    when: result|failed
  - command: /bin/something_else
    when: result|success
  - command: /bin/still/something_else
    when: result|skipped
```

Overview

1. Simplicity in mind
2. Modules

```
$ pwd
/Users/jdewey/git/ansible/library
$ ls cloud/* | wc -l
62
$ find . -type f | wc -l
229
```

Overview

1. Simplicity in mind
2. Modules

```
1 ---
2 - name: provision test instances
3   hosts: local
4   connection: local
5   vars_files:
6   - ../vars/main.yml
7   tasks:
8   - include: keypair.yml
9   - name: create the security group and rules
10     nova_group:
11       name: "{{ testenv_security_groups }}"
12       description: "{{ testenv_security_groups_description }}"
13       rules:
14         - ip_protocol: "{{ item.proto }}"
15           from_port: "{{ item.port }}"
16           to_port: "{{ item.port }}"
17           cidr: 0.0.0.0/0
18           state: "{{ item.state }}"
19       with_items: testenv_security_group_rules
20       register: testenv_security_group
21     - command: neutron security-group-rule-create --ethertype={{ item }} --remote-group-id={{ testenv_security_group.results[0].group_id }}
22       register: security_group_rule_create_result
23       failed_when: "security_group_rule_create_result.rc != 0 and 'Security group rule already exists.' not in security_group_rule_create_result.stderr"
24       changed_when: security_group_rule_create_result.rc == 0
25     with_items:
26     - IPv4
27     - IPv6
28   - name: create {{ item }}
29     nova_compute:
30       name: "{{ item }}"
31       image_id: "{{ testenv_image_id }}"
32       key_name: "{{ testenv_keypair_name }}"
33       security_groups: "{{ testenv_security_group.results[0].group_id }}"
34       wait_for: 200
35       flavor_id: 3
36       nics:
37         - net-id: "{{ testenv_net_id }}"
38     with_items: testenv_instance_names
39   - name: associate a floating IP to {{ item }}
40     nova_fip: server={{ item }}
41     with_items: testenv_instance_names
42     register: testenv_floating_ips
43   - name: wait for {{ item }} to boot
44     wait_for: port=22 delay=5 timeout=300 host={{ item.floating_ip }}
45   with_items: testenv_floating_ips.results
```

```
graph TD
    18 --> 20
    37 --> 41
    37 --> 43
```

Overview

1. Simplicity in mind
2. Modules

```
89 def main():
90     module = AnsibleModule(
91         argument_spec = dict(
92             login_username = dict(default='admin'),
93             login_password = dict(required=True),
94             login_tenant_name = dict(required=True),
95             auth_url = dict(default='http://127.0.0.1:35357/v2.0/'),
96             region_name = dict(default=None),
97             name = dict(required=True),
98             public_key = dict(default=None),
99             state = dict(default='present', choices=['absent', 'present'])
100         ),
101     )
102
103     nova = nova_client.Client(module.params['login_username'],
104                               module.params['login_password'],
105                               module.params['login_tenant_name'],
106                               module.params['auth_url'],
107                               service_type='compute')
108
109     try:
110         nova.authenticate()
111     except exc.Unauthorized, e:
112         module.fail_json(msg = "Invalid OpenStack Nova credentials.: %s" % e.message)
113     except exc.AuthorizationFailure, e:
114         module.fail_json(msg = "Unable to authorize user: %s" % e.message)
115
116     if module.params['state'] == 'present':
117         for key in nova.keypairs.list():
118             if key.name == module.params['name']:
119                 module.exit_json(changed = False, result = "Key present")
120
121         try:
122             key = nova.keypairs.create(module.params['name'], module.params['public_key'])
123         except Exception, e:
124             module.exit_json(msg = "Error in creating the keypair: %s" % e.message)
125         if not module.params['public_key']:
126             module.exit_json(changed = True, key = key.private_key)
127         module.exit_json(changed = True, key = None)
128     if module.params['state'] == 'absent':
129         for key in nova.keypairs.list():
130             if key.name == module.params['name']:
131                 try:
132                     nova.keypairs.delete(module.params['name'])
133                 except Exception, e:
134                     module.fail_json(msg = "The keypair deletion has failed: %s" % e.message)
135                 module.exit_json( changed = True, result = "deleted")
136             module.exit_json(changed = False, result = "not present")
137
138 # this is magic, see lib/ansible/module.params['common.py']
139 from ansible.module_utils.basic import *
```

Overview

1. Simplicity in mind
2. Modules
3. Inventory

Hosts

[identity] ← Group
identity-1.example.com
identity-2.example.com

Overview

1. Simplicity in mind
2. Modules
3. Inventory

Variables

identity:

port: 5000

admin_port: 35357

auth_strategy: uuid


...

Overview

1. Simplicity in mind
2. Modules
3. Inventory
4. Playbooks

```
---  
- hosts: identity  
  roles:  
    - base  
    - memcached  
    - apache2  
    - { role: keystone,  
identity.auth_strategy: pki }
```

Group Targeting



Overview

1. Simplicity in mind
 2. Modules
 3. Inventory
 4. Playbooks
 5. Roles
- Roles are tasks
 - Roles are to be included in playbooks
 - Roles are shareable
<https://galaxy.ansible.com>

Overview

1. Simplicity in mind
2. Modules
3. Inventory
4. Playbooks
5. Roles

```
roles/  
  identity/  
    tasks/  
    handlers/  
    files/  
    templates/  
    defaults/  
    vars/  
    meta/
```

Overview

1. Simplicity in mind
2. Modules
3. Inventory
4. Playbooks
5. Roles
6. You already have everything you need



Orchestration

```
31 - name: install percona arbiter
32 hosts: db_arbiter
33 roles:
34   - percona-common
35   - percona-arbiter
36
37 - name: configure percona
38 hosts: db[0]
39 roles:
40   - percona-arbiter
41
42 - name:
43 hosts:
44 roles:
45   - ne
46
47 - name:
48 hosts:
49 roles:
50   - ke
51
52 - name:
53 hosts:
54 vars_f
55 - ro
56 roles:
57   - ne
58
59 - hosts:
60 roles:
61   - op
62
63 - name:
64 hosts:
65 roles:
66   - ne
67
68 - name:
69 hosts:
70 tasks:
71   - se
72
73 - name:
74 hosts:
75 vars_f
76 - ro
77 roles:
78   - no
79
80 - name:
81 hosts:
82 roles:
83   - sw
84
85 - name:
86 hosts:
87 roles:
88   - sw
89
90 - name:
91 hosts:
92 roles:
93   - ha
94   - sw
95   - sw
96   - sw
97   - sw
98
99 - name:
100 hosts:
101 roles:
102   - gl
103
104 - name:
105 hosts:
106 vars_f
107 - ro
108 roles:
109   - ci
110
111 - name:
112 hosts:
113 roles:
114   - ha
115
116 - name:
117 hosts:
118 roles:
119   - db
120
121 - name:
122 hosts:
123 roles:
124   - ci
125
126 - name:
127 hosts:
128 roles:
129   - ho
130   - ha
131   - co
132
133 - hosts:
134 roles:
135   - co
136   - ci
137
138 - hosts:
139 roles:
140   - op
141
142 - name: restart neutron-openvswitch-agent
143 hosts: network:compute
144 tasks:
145   - service: name=neutron-openvswitch-agent
146
147 - name: Install monitoring
148 include: playbooks/monitoring.yml
```

```
80 - name: swift common code and config
81   hosts: swiftnode
82   roles:
83     - swift-common
84
85 - name: swift bootstrap rings
86   hosts: swiftnode_primary
87   roles:
88     - swift-ring
89
90 - name: swift code and config
91   hosts: swiftnode
92   roles:
93     - haproxy
94     - swift-object
95     - swift-account
96     - swift-container
97     - swift-proxy
98
99 - name: glance code and config
100   hosts: controller:db
101   roles:
102     - glance-common
103
```

```
39 78 - nova-common
40 79
```

Orchestration

branch: **master** ▾

ansible-swift-private-cloud / roles / proxy-server / tasks / **main.yml**



 3 months ago Pass -E environment into chef-client run

1 contributor

file | 6 lines (4 sloc) | 0.235 kb



Open

Edit

Raw

Blame

History

Delete

```
1  #- name: Drop Client JSON
2      #template: src=./etc/chef/chef-client.json.j2 dest=/etc/chef/chef-client.json
3
4  - name: Chef Proxy Role
5      shell: chef-client -E swift-private-cloud -o "role[spc-starter-proxy]" -L "/etc/chef/chef-client.log"
```

Why?

```
@roles('mgmt')
def _create_default_network():
    with cd('/root'):
        run('. openrc && neutron net-create flat1004 --provider:network_type flat
        --provider:physical_network physnet1')
```

- Not idempotent
- Everyone is rolling their own

When

```
- quantum_network: >  
  name=flat1004 state=present  
  provider_network_type=flat  
  provider_physical_network=physnet1'
```

One Time Tasks

```
- name: cleanup expired keystone database tokens  
  shell: mysql -e "delete from keystone.token where expires < date_sub(now(),  
interval 24 hours);"
```

One Time Tasks

```
---  
- name: is m1.tiny undersized?  
  shell: mysql -e "select root_gb from nova.instance_types where name='m1.tiny';" | grep 10  
  ignore_errors: True  
  changed_when: False  
  register: resize_tiny_flavor  
  
- name: bump root disk size on m1.tiny  
  shell: mysql -e "update nova.instance_types set root_gb=10 where name='m1.tiny';"  
  when: resize_tiny_flavor.rc != 0
```



Operations

```
1 ---
2 - hosts: all
3 vars:
4   openssl_packages: ["openssl", "libssl1.0.0"]
5   openssl_impacted_service:
6     - nginx
7     - apache2
8     - postgresql
9     - php5-fpm
10    - openvpn
11    - postfix
12    - monit
13    - zabbix-server
14 tasks:
15   - name: ensure openssl is the last version
16     apt: pkg={{item}} state=latest update_cache=yes
17     register: openssl_updated
18     with_items: openssl_packages
19     when: ansible_os_family == "Debian"
20
21   - name: check if service need to be restarted
22     shell: "lsof -n | grep 'DEL.*libssl.so'"
23     register: result_check
24     failed_when: result_check.rc > 1
25     changed_when: result_check.rc != 1
26     always_run: yes
27
28   - name: test running services
29     command: "service {{item}} status | grep -i running"
30     register: services_status
31     with_items: openssl_impacted_service
32     when: result_check.rc == 0 or openssl_updated.changed
33     ignore_errors: true
34     always_run: yes
35
36   - name: restart running service
37     service: name={{item.item}} state=restarted
38     with_items: services_status.results
39     when: (result_check.rc == 0 or openssl_updated.changed ) and item.rc == 0
40
41   - name: ensure no more service need to be restarted
42     shell: "lsof -n | grep 'DEL.*libssl.so'"
43     register: result
44     failed_when: result.rc == 0
45     changed_when: result.rc != 1
46     always_run: yes
```

Integration Testing

```
1 ---
2 - hosts: controller[0]
3   tasks:
4     - name: migrate neutron services to test-controller-0
5       shell: . /root/stackrc; HOSTNAME=test-controller-0 /usr/local/bin/migrate_neutron_services
6     - name: neutron agents are all alive
7       shell: . /root/stackrc; neutron agent-list | awk '/ xxx / {print;ec=1} END{exit ec}'
8     - name: neutron has an internal network
9       shell: . /root/stackrc; neutron net-list | grep internal
10    - name: neutron has a network with network_type vxlan
11      shell: . /root/stackrc; neutron net-show internal | grep provider:network_type | grep vxlan
12    - name: neutron has a network with segmentation_id 256
13      shell: . /root/stackrc; neutron net-show internal | grep provider:segmentation_id | grep 256
14    - name: neutron has a network with router_external False
15      shell: . /root/stackrc; neutron net-show internal | grep router:external | grep False
16    - name: neutron has a network with internal_subnet
17      shell: . /root/stackrc; neutron net-list | grep internal | grep 172.16.255.0/24
18    - name: neutron has the internal subnet
19      shell: . /root/stackrc; neutron subnet-list | grep internal
20    - name: neutron has the internal subnet with cidr
21      shell: . /root/stackrc; neutron subnet-show internal | grep cidr | grep 172.16.255.0/24
22    - name: neutron has the internal_subnet with cidr start/end addresses
23      shell: . /root/stackrc; neutron subnet-show internal | grep allocation_pools | egrep '172.16.255.2.*172.16.255.254'
24    - name: neutron has the internal_subnet with enable_dhcp True
25      shell: . /root/stackrc; neutron subnet-show internal | grep enable_dhcp | grep True
26    - name: neutron has the internal_subnet with gateway_ip
27      shell: . /root/stackrc; neutron subnet-show internal | grep gateway | grep gateway_ip
28    - name: neutron has the default router
29      shell: . /root/stackrc; neutron router-list | grep default
30    - name: neutron router can ping internet
31      shell: ROUTER_NS=$( ip netns show | grep qrouter- ); ip netns exec ${ROUTER_NS} ping -c 5 8.8.8.8
32
33 - hosts: controller
34   tasks:
35     - name: neutron dnsmasq has 8.8.8.8 upstream resolver
36       shell: grep 8.8.8.8 /etc/dnsmasq.conf
37     - name: neutron dnsmasq has 8.8.4.4 upstream resolver
38       shell: grep 8.8.8.8 /etc/dnsmasq.conf
39     - name: neutron config has rabbit servers
40       shell: egrep "rabbit_hosts = [0-9.]+:5672,[0-9.]+:" /etc/neutron/neutron.conf
41     - name: iptables mangle rule in place to correct DHCP checksums
42       shell: iptables -L -n -t mangle | egrep '^CHECKSUM\s+udp\s+--\s+0.0.0/0\s+0.0.0/0\s+udp dpt:68 CHECKSUM fill'
```

Integration Testing

```
1 # UNINSTALL
2 - name: uninstall hello with apt
3   apt: pkg=hello state=absent purge=yes
4   register: apt_result
5
6 - name: check hello with dpkg
7   shell: dpkg --get-selections | fgrep hello
8   failed_when: False
9   register: dpkg_result
10
11 - debug: var=apt_result
12 - debug: var=dpkg_result
13
14 - name: verify uninstallation of hello
15   assert:
16     that:
17       - "'changed' in apt_result"
18       - "dpkg_result.rc == 1"
19
20 # UNINSTALL AGAIN
21 - name: uninstall hello with apt
22   apt: pkg=hello state=absent purge=yes
23   register: apt_result
24
25 - name: verify no change on re-uninstall
26   assert:
27     that:
28       - "not apt_result.changed"
29
30 # INSTALL
31 - name: install hello with apt
32   apt: name=hello state=present
33   register: apt_result
34
35 - name: check hello with dpkg
36   shell: dpkg --get-selections | fgrep hello
37   failed_when: False
38   register: dpkg_result
39
40 - debug: var=apt_result
41 - debug: var=dpkg_result
42
43 - name: verify installation of hello
44   assert:
45     that:
46       - "apt_result.changed"
47       - "dpkg_result.rc == 0"
```

Common patterns

```
22 {% macro rabbitmq_hosts() -%}  
23 {% for host in groups['controller'] -%}  
24     {% if loop.last -%}  
25     {{ hostvars[host][primary_interface]['ipv4']['address'] }}:{{ rabbitmq.port }}  
26     {%- else -%}  
27     {{ hostvars[host][primary_interface]['ipv4']['address'] }}:{{ rabbitmq.port }},  
28     {%- endif -%}  
29 {% endfor -%}  
30 {% endmacro -%}  
31  
32 {% if rabbitmq.cluster -%}  
33 rabbit_hosts = {{ rabbitmq_hosts() }}  
34 {% else -%}  
35 rabbit_host = {{ endpoints.rabbit }}  
36 rabbit_port = 5672  
37 {% endif -%}
```

```
$ grep -R 'macro rabbitmq_hosts' . | wc -l
```

Variables can be tricky

```
$ cat memcached/defaults/main.yml
```

```
memcached:
```

```
  port 11211
```

```
# from site.yml
```

```
- name: openstack horizon service
```

```
  hosts: controller
```

```
  roles:
```

```
    - memcached # must be added to roles
```

```
    - horizon   # for horizon to reference {{ memcached.port }}
```


Variables can be tricky

```
$ cat group_vars/all.yml
```

```
rabbitmq:
```

```
  cluster: true
```

```
# from site.yml
```

```
- name: nova code and config
```

```
  hosts: controller
```

```
  vars_files:
```

```
    # don't do this -- will reset the override from inventory
```

```
    - roles/rabbitmq/defaults/main.yml
```

```
  roles:
```

```
    - nova-common
```

Variables can be tricky

We ended up defining everything in inventory, and managing with a custom vars plugin.

```
$ grep vars_plugin ansible.cfg  
vars_plugins = plugins/vars
```



References

- Ansible: <https://github.com/ansible/ansible>
- Ansible Docs: <http://docs.ansible.com>
- Ansible Swift: <https://github.com/rcbops/ansible-swift-private-cloud>
- Heartbleed: <https://github.com/jdauphant/patch-openssl-CVE-2014-0160>
- Ursula: <https://github.com/blueboxgroup/ursula/>

Success.

